

# E 0243: Assignment 1

Aman Sachan (18094)

Paras Lohani (18226)

## 1 Introduction

CPI is cycle per instruction which symbolizes the number of cycles taken by an instruction. It depends on the cycles taken by the various hardware events. Each CPI contains a base CPI which reflects the minimum CPI needed to execute instructions. Since we are experimenting in a superscalar processor, there is a chance of overlapping of hardware events which will effect CPI accordingly. We will run experiment on various benchmarks to understand the performance of our processor and the impact of various hardware events on it.

## 2 Hardware Performance Counters

Hardware Performance Counters is a type of registers which stores the count of hardware related events performed in the computer systems.

We have done all the computations in the Intel(R) Core(TM) i3-6006U CPU @ 2.00GHz processor for the following SPECINT Benchmarks 531.deepsjeng\_r, 520.omnetpp\_r, 557.xz\_r and for SPEC FP Benchmarks 526.blender\_r, 507.cactuBSSN\_r, 511.povray\_r.

### 2.1 Performace monitoring tool

Open source tool named **perf** has been used to collect supported hardware events counters on a given program. It gives facility of command line interface where one can use commands to collect, analyze performance and trace data. The **perf** commands used are **stat**, **record**, **report**, etc.

### 2.2 Data Collection

Using **perf** we collected the measures of performance counters in the user space. The data is collected for atleast these following 11 events:-

- |                  |                         |                      |
|------------------|-------------------------|----------------------|
| • cycles         | • L2-store-misses       | • dTLB-store-misses  |
| • instructions   | • L1-dcache-load-misses | • iTLB-load-misses   |
| • branch-misses  | • L1-icache-load-misses |                      |
| • L2-load-misses | • dTLB-load-misses      | • branch-load-misses |

We have used the command

```
perf stat -I interval(in ms) -e cycles:u,instructions:u,branch-misses:u,  
L1-dcache-load-misses:u,L1-icache-load-misses:u,dTLB-load-misses:u,dTLB-  
store-misses:u,iTLB-load-misses:u,branch-load-misses:u,L2-load-misses:u,  
L2-store-misses:u bash ./run.sh taskset--cpu-list 0 2>filename.txt
```

to generate the dataset in the file and set the benchmark to run on only one core of the processor after that we have to pre-process it(data cleaning), to get data in proper format.

### 2.3 Data Cleaning

Data collected from the **perf** tool is in row format where each row depicts the counter information for a particular hardware events. Data in a row is in the form of [Time, Counter value, Event name]. These hardware events will be used as independent variables in our model, therefore pre-processing needs to be done on the data collected from the previous step to transform it to column format. The counters collected are comma separated which is not suitable for model creation, therefore some cleaning steps need to be performed on the dataset. While analysing values of columns we have observed some outliers which is removed. This outliers mainly has CPI way more than the mean CPI collected.

### 2.4 Data Visualization

Figure 1. shows the scatter plot of various events against CPI in the 557.xz\_r benchmark. We can observe that there is a positive slope linear relationship of each events with the CPI. This conclude that all the events contribute in an additive fashion.

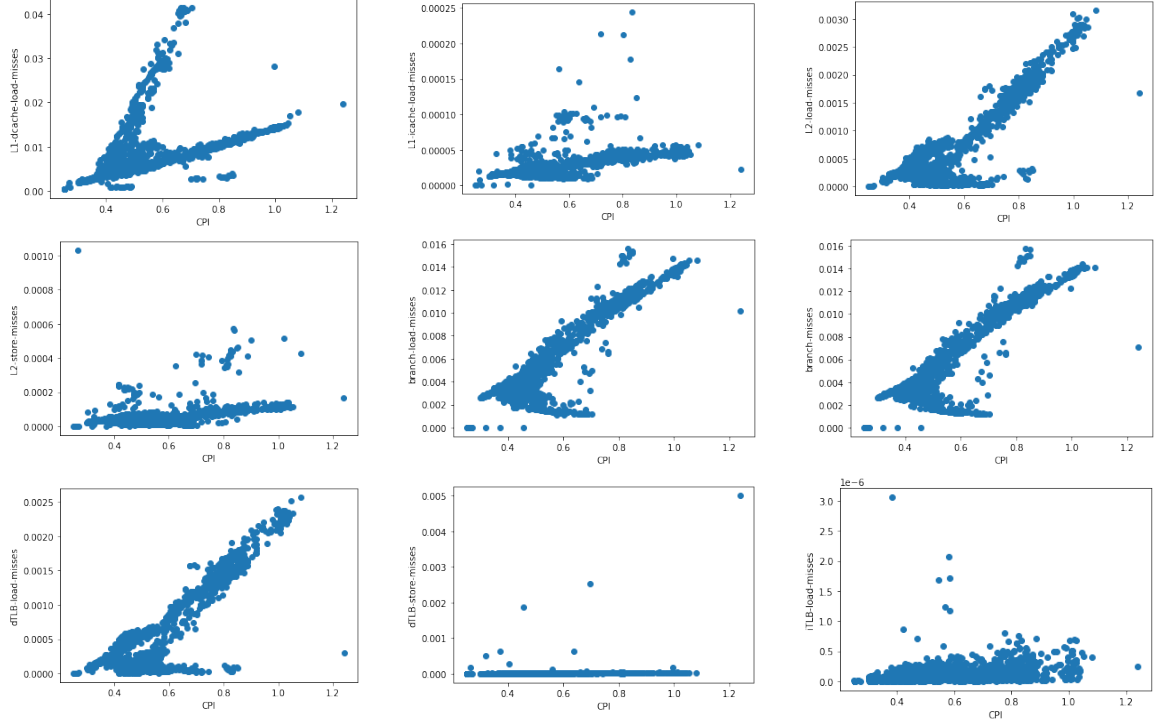


Figure 1: Various hardware events against CPI

## 2.5 Basic Statistical Measures

After cleaning data, we analyzed the distribution of data column wise. For each column we get the information like count, mean, std, min, max etc.

```
df.describe()
```

	L1-icache-load-misses	L2-load-misses	instructions	branch-misses	iTLB-load-misses	L1-dcache-load-misses	branch-load-misses	L2-store-misses	cycles	dTLB-load-misses	dTLB-store-misses	cpi
count	1577.000000	1.577000e+03	1.577000e+03	1.577000e+03	1577.000000	1.577000e+03	1.577000e+03	1.577000e+03	1.577000e+03	1.577000e+03	1.577000e+03	1577.000000
mean	26198.035510	5.858331e+05	1.035216e+09	5.066215e+06	120.257451	9.421840e+06	5.065176e+06	5.425729e+04	5.449614e+08	4.840591e+05	1.700226e+04	0.606550
std	14051.350822	4.305526e+05	3.022216e+08	2.117000e+06	158.276223	7.133579e+06	2.119018e+06	5.239267e+04	2.165339e+07	3.761005e+05	7.202643e+04	0.995811
min	764.000000	6.200000e+02	2.234884e+06	0.000000e+00	0.000000	1.591000e+04	0.000000e+00	0.000000e+00	7.923140e+07	5.800000e+01	0.000000e+00	0.251153
25%	20544.000000	2.164470e+05	8.033458e+08	3.846675e+06	51.000000	5.590731e+06	3.815301e+06	3.007800e+04	5.466393e+08	1.545510e+05	6.409000e+03	0.437905
50%	25829.000000	5.066450e+05	1.040864e+09	5.159275e+06	95.000000	6.895652e+06	5.167480e+06	4.910100e+04	5.469279e+08	3.813130e+05	1.268500e+04	0.525590
75%	29154.000000	8.720860e+05	1.245433e+09	7.001565e+06	156.000000	9.285867e+06	7.020379e+06	6.516700e+04	5.471902e+08	7.957090e+05	1.563100e+04	0.679146
max	158379.000000	1.622668e+06	2.163339e+09	1.007493e+07	4150.000000	3.405189e+07	1.015340e+07	1.093789e+06	5.497796e+08	1.315614e+06	1.413938e+06	35.452133

## 2.6 Normalization

Instead of normalizing the data using traditional min-max, range-transformation etc., we are dividing the counters of various events with the corresponding instruction counter. This way we will deal with 0-1 data values.

## 3 Learning Model using Linear Regression

### 3.1 Sampling

As different benchmarks take different amount of running time. So, we have taken different intervals according to the running time of the benchmark. We have computed approx 1250+ datasets for each benchmark and trained our Linear Regression model on random 80% of the datasets (1000+) and test it on the remaining 20% of the datasets (250+) on average.

### 3.2 Linear Regression

We have used a Multiple Linear Regression model and a base as Linear Regression. To build the model we used `python` and some already available python libraries such as `numpy`, `matplotlib`, `statsmodels`, `pandas`, `sklearn`, `math`, `seaborn`.

To analyze the CPI stack i.e, how the features are impacting the CPI stack, we restricted our coefficients to be non-negative.

### 3.3 Error Analysis

We have used different error analyzing factor such as `R2`, `Residuals`, `F-statistic`, `p-value`, `RMSE` and `Adjusted R2` values to assess the quality of the model generated.

## 4 Code

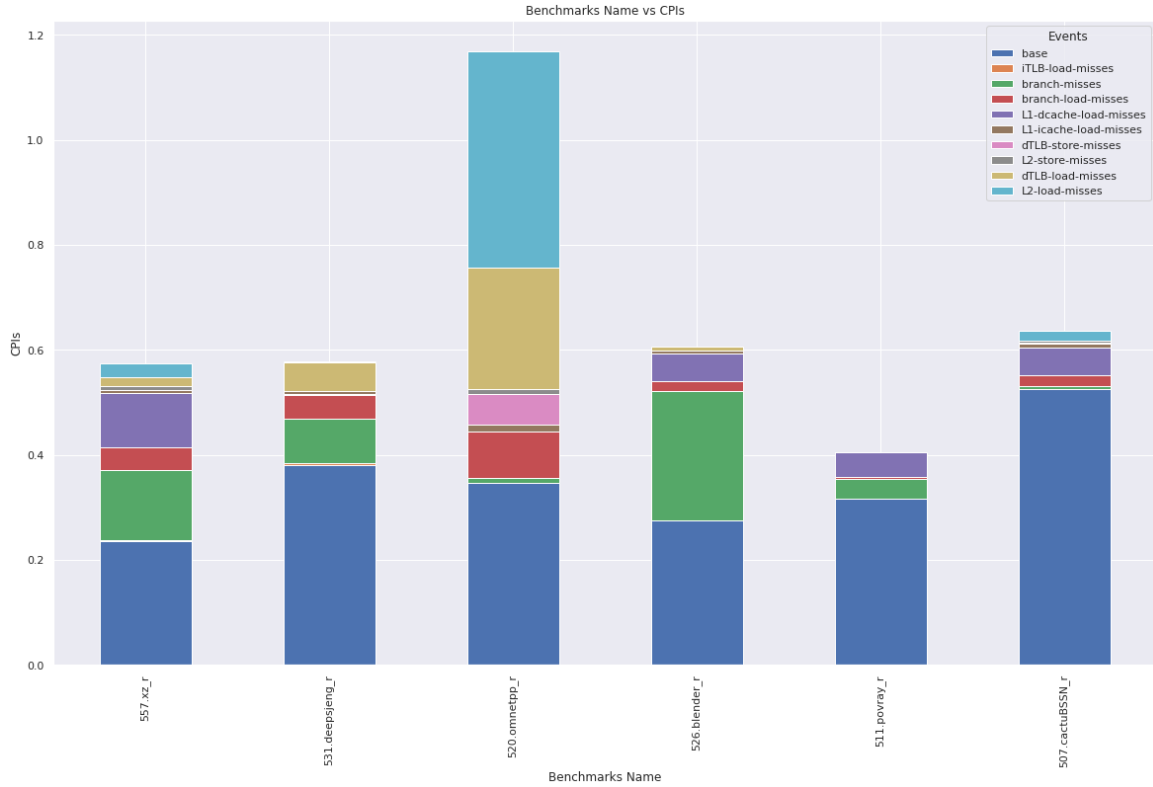
Code is written in **Python** using **Google Colab**. For the preprocessing, one can refer [this link](#). For the linear regression models, one can refer [this link](#). Final code having all the code that can be run on all the benchmarks for the cpi stack creation.

## 5 Results

CPI Stack(dependent variable) is computed on our test hardware events using our trained Linear Regression model. This model uses the equation :-

$$y = \beta_0 + \beta_1 * x_1 + \beta_2 * x_2 + \dots$$

where,  $y$  will be estimated CPI,  $\beta_0$  is base CPI,  $\beta_i$  is the coefficient corresponding to the feature  $x_i$ . Coefficients is representing the number of cycles required for a particular miss event and the factors are in the term of number of miss events per instruction.



## 5.1 SPECINT Benchmarks

### 5.1.1 557.xz\_r

Number of dataset are 1575 and we trained on random 80% of the dataset. We have taken the interval to be 275 ms and run our model on whole interval.

events	mean	coefficients	p-value
L1-icache-load-misses	2.8024535566934098e-05	180.78003596543587	0.000
dTLB-load-misses	0.0005953550353532326	32.26056892641485	0.009
iTLB-load-misses	1.2934668363220614e-07	8796.132954040599	0.003
L1-dcache-load-misses	0.010080777594353064	10.279066147721124	0.000
L2-load-misses	0.0007150048170411444	32.756427934147226	0.390
L2-store-misses	5.9448583747928424e-05	124.94326595805428	0.020
branch-load-misses	0.005702752243598566	4.02678437371891	0.151
dTLB-store-misses	2.01680480267406e-05	90.3132679632431	0.000
branch-misses	0.005703521745964653	26.819936622290964	0.000

RMSE : 0.016767040724761286, R2 : 0.9886135941397272, F-value : 2450.3669452945906, Adjusted R2 value : 0.9885348257671349, Residuals(abs mean) : 0.012054561925430718

Intercept (base CPI) : 0.2345000696782411

original CPI : 0.5734389757747521

estimated CPI : 0.5735778596039196

**Observation :** base CPI is contributing the most CPI after that branch-misses and after that L1-dcache-load-misses and others are some part of the overall CPI, this means there are large number of branch instructions.

### 5.1.2 531.deepsjeng\_r

Number of dataset are 1571 and we trained on random 80% of the dataset. We have taken the interval to be 350 ms and run our model on whole interval.

events	mean	coefficient	p-value
L1-icache-load-misses	0.0007591502606098287	10.548834608709802	0.000
dTLB-load-misses	0.00018243457434705758	260.09942954806377	0.000
iTLB-load-misses	1.9359910563306549e-07	16319.28074338875	0.000
L1-dcache-load-misses	0.002763126731329056	1.4155543849647911	0.000
L2-load-misses	0.00032576283515684126	23.567211520135192	0.002
L2-store-misses	8.903333087539272e-05	3.1918324013308803	0.000
branch-load-misses	0.006053893386710748	7.1739947447410275	0.002
dTLB-store-misses	2.3894614896774562e-06	211.1452901664822	0.000
branch-misses	0.006052024093601827	14.11053394324681	0.000

RMSE : 0.0044205014775149185, R2 : 0.945488246957873, F-value : 487.5778833225608, Adjusted R2 value : 0.9451102764051926, Residuals (abs mean) : 0.0033498693068257568

Intercept (base CPI) : 0.3803417648984122

original CPI : 0.5779147230870468

estimated CPI : 0.5779257701748101

**Observation :** base CPI is contributing more than 50% of the CPI after that branch-misses and after that dTLB-load-misses and after that branch-load-misses and others are some part of the overall CPI.

### 5.1.3 520.omnetpp\_r

Number of dataset are 1421 and we trained on random 80% of the dataset. We have taken the interval to be 450 ms and run our model on whole interval.

events	mean	coefficient	p-value
L1-icache-load-misses	0.0010798950442872002	14.780335417438875	0.000
dTLB-load-misses	0.004758378019042945	54.09271037780789	0.000
iTLB-load-misses	4.3589142753375e-06	0	0.000
L1-dcache-load-misses	0.03744935097402198	0	0.000
L2-load-misses	0.0074454754789509345	53.474179409831315	0.000
L2-store-misses	0.001046928816870187	39.188127003564006	0.886
branch-load-misses	0.004710299807015788	17.395182607818338	0.465
dTLB-store-misses	0.0005545226762658645	0	0.000
branch-misses	0.004709991278449441	9.719290929948073	0.024

RMSE : 0.003990603427194033, R2 : 0.9848162655220875, F-statistic : 1643.11874413056, Adjusted R2 value : 0.984699766280569, Residuals (abs mean) : 0.0028522254443616203

Intercept (base CPI) : 0.32666649256775715  
original CPI : 1.1685336335622385  
estimated CPI : 1.168508292658236

**Observation :** L2-load misses is contributing most of the CPI stack after that base CPI and after that dTLB-load-misses and others are some part of the overall CPI, this means many instructions are missed during L2 cache load.

## 5.2 SPEC FP Benchmarks

### 5.2.1 526.blender\_r

Number of dataset are 1468 and we trained on random 80% of the dataset. We have taken the interval to be 350 ms and run our model on whole interval.

events	mean	coefficient	p-value
L1-icache-load-misses	0.0007699121439311228	7.388860482313896	0.000
dTLB-load-misses	0.00014065278675612873	45.95655182791539	0.000
iTLB-load-misses	8.705246159453715e-07	0	0.000
L1-dcache-load-misses	0.0077514567943604415	7.016206644356411	0.000
L2-load-misses	0.0004467760649263368	0	0.000
L2-store-misses	3.422228108340457e-05	0	0.000
branch-load-misses	0.005887557188933684	4.795891812372351	0.002
dTLB-store-misses	2.8367396025371487e-06	76.72813571118569	0.000
branch-misses	0.005896845768060546	39.63353414474996	0.000

RMSE : 0.02861564792853203, R2 : 0.8810636209931347, F-value : 194.25045771961976, Adjusted R2 value : 0.8801804300599154, Residuals : 0.019904452932037697

Intercept (base CPI) : 0.2650321312941003  
original CPI : 0.6061072876866389  
estimated CPI : 0.6063202794881527

**Observation :** base CPI and branch-misses are contributing the most part of the CPI stack after that L1-dcache-load-misses and others are some part of the overall CPI, this means many branch instructions are miss-predicted.

### 5.2.2 511.povray\_r

Number of dataset are 1264 and we trained on random 80% of the dataset. We have taken the interval to be 600 ms and run our model on whole interval.

events	mean	coefficient	p-value
L1-icache-load-misses	0.0029576525932153044	0.0	0.000
dTLB-load-misses	2.1163698394802664e-07	0.0	0.000
iTLB-load-misses	5.898601446829382e-07	0.0	0.000
L1-dcache-load-misses	0.021675541848364625	2.169353937225385	0.000
L2-load-misses	2.4891987772304925e-07	15.086639719250178	0.739
L2-store-misses	2.578724500302794e-08	0.0	0.000
branch-load-misses	0.0011507021665749242	3.6756199081719525	0.022
L1-icache-load-misses	4.926980361377182e-09	0.0	0.000
branch-misses	0.0011525716937059072	32.57901986937401	0.168

RMSE : 0.0018729584255113394, R2 : 0.9713923963675041, F-value : 762.1177556186664, Adjusted R2 value : 0.9711453057411198, Residuals (abs mean) : 0.0012857335832263746

Intercept (base CPI) : 0.31628257708657365  
original CPI : 0.4052723621588838  
estimated CPI : 0.4052359546588299

**Observation :** base CPI is contributing more than 50% of the CPI stack after that branch-misses and L1-dcache-load-misses and others are some part of the overall CPI.

Most of the coefficients are 0 that means that particular events are not contributing anything in the CPI Stack this also means that these events are executed in parallel with other instructions.

### 5.2.3 507.cactuBSSN\_r

Number of dataset are 1203 and we trained on random 83.3% of the dataset. We have taken the interval to be 350 ms and run our model on whole interval.

events	mean	coefficient	p-value
L1-icache-load-misses	0.011551163926960761	0.2760627552846979	0.000
dTLB-load-misses	9.507220557099351e-05	0.0	0.000
iTLB-load-misses	5.385365294963261e-07	0.0	0.000
L1-dcache-load-misses	0.10055920061624948	0.8193610224462698	0.000
L2-load-misses	0.002283097493094351	4.891398632062899	0.000
L2-store-misses	0.000827177658804286	5.489462627476771	0.836
branch-load-misses	1.3942384804888686e-05	1190.1643967253897	0.000
dTLB-store-misses	2.6987977046142073e-05	112.5820392384979	0.000
branch-misses	1.3848659400287552e-05	45.60084100731926	0.000

RMSE : 0.018881266672916255, R2 : 0.725912320285156, F-value : 56.20636648450638, Adjusted R2 value : 0.7234256377070979, Residuals (abs mean) : 0.011893709370990749

Intercept (base CPI) : 0.5130003480309733  
original CPI : 0.6354518768130953  
estimated CPI : 0.6354242303992013

**Observation :** base CPI is contributing more than 50% of the CPI stack after that L1-dcache-load-misses and others are some part of the overall CPI.

## 6 Conclusion

In 5 out of 6 benchmarks, we are getting  $CPI < 1$  which shows our processor is superscalar and in some benchmarks, some coefficients are 0.0 which shows the instructions are overlapping with other instructions. Also in one of the benchmark 520.omnetpp\_r we are getting  $CPI > 1$  which means that by using superscalar processor it is not guarantee that we will always get more than 1 instruction per cycle.